

Technical report for a DEVS formalism simulator in C++

Bc. Jozef Méry

xmeryj00@vut.cz

ABSTRACT

This technical report is concerned with a *DEVS formalism simulator* written in *C++* as a part of the project assignment of subject *Simulation Tools and Techniques (SNT)* in 2023. The report presents the *formal definitions* of the *atomic*, and *compound* DEVS models. Next, the report shares some details about the *library implementation*. Finally, a *demo application*, and its *examples* are discussed.

1 Introduction

The aim of the project was to create a *DEVS formalism simulator* in *C++* using the standard library only, and provide an example from the *queue theory* area. This technical report introduces the *DEVS formalism* in this section, describes the *simulator implementation* in section 2, discusses the demo application and its examples in section 3, and finally the report is concluded in section 4.

The DEVS (*Discrete Event System Specification*) formalism is a structure defined by Ziegler et al. [2] for modeling and analyzing system with discrete events. There are two models in *DEVS*, namely *atomic*, and *compound* models. The *atomic model* is the smallest standalone unit in *DEVS*, and is defined as follows:

$$M = (X, Y, S, \delta_{int}, \delta_{ext}, \lambda, ta), \quad (1)$$

where X is the set of *input* values, Y is the set of *output* values, and S is the set of *states*. These sets are often *structured* as $X = \{(x_1, x_2, \dots, x_m) | x_1 \in X_1, x_2 \in X_2, \dots, x_m \in X_m\}$. Individual members in the structured input and output sets are referred to as *ports* (the provided example has m ports). Next, δ_{int} is the *internal transition function* ($\delta_{int} : S \rightarrow S$), δ_{ext} is the *external transition function* ($\delta_{ext} : Q \times X \rightarrow S$, where Q is the *total state* defined as $Q = \{(s, e) : s \in S, 0 \leq e \leq ta(s)\}$), λ is the *output function* ($\lambda : S \rightarrow Y$), and finally ta is the *time advance function* ($ta : S \rightarrow \mathbb{R}_0^+ \cup \{\infty\}$), defining the *remaining* time until the next internal transition. [1][2]

The *compound DEVS* model is defined as follows:

$$N = (X, Y, D, \{M_d | d \in D\}, \{I_d | d \in D \cup \{N\}\}, \{Z_{i,d} | d \in D \cup \{N\}, i \in I_d\}, Select), \quad (2)$$

where X is again the set of *input* values, Y is the set of *output* values, D is the set of *atomic DEVS component references* (names), $M_d | d \in D$ is the set of *atomic DEVS components* (see definition 1), $I_d | d \in D \cup \{N\}$ is the set of *influencers* without a self-loop ($d : I_d \subseteq D \cup \{N\}, d \notin I_d$), $Z_{i,d} | i, d \in D \cup \{N\}$ is the set of *influencer transform functions* defining *connections* between components. Components can be *connected* to any other components, or the enclosing compound model N according to the following *rules*:

$$Z_{N,d} : X \rightarrow X_d | d \in D, \quad Z_{i,N} : Y_i \rightarrow Y | i \in I_N, \quad Z_{i,d} : Y_i \rightarrow X_d | d \in D, i \in I_d, \quad (3)$$

implying that the components cannot directly *send inputs* to, or *receive outputs* from the compound model N . Instead, they contribute to the *compound input* X and *compound output* Y . Finally, *Select* is a *tie-breaking function* ($2^D \rightarrow D$) in case of *concurrent events*. Additionally, it is *formally proven* that every *compound DEVS model* N is an *atomic DEVS model* at the same time. Thus components can be *arbitrarily nested*, allowing the modeling of *hierarchical systems*. [1][2]

2 Library description

The library is written in C++17, and is provided as a *single-file, header-only* library for easier potential installation. It implements the fundamental elements of DEVS, namely *atomic models*, *compound models*, and a *simulator* based on the *next-event (calendar)* algorithm in an *object-oriented* fashion. Additional features include *printing*, *simplified random number generation*, *helper constants*, or *helper types*. Any and all library code is located in the *Devs* name space.

The core element of the *simulator* is a calendar structure using the *next-event algorithm*. The implementation utilizes the *priority queue* from the standard library specialized with a custom *Event* structure. The calendar itself has only *two* levels of *priority hierarchy*, namely *remaining time*, and then *FIFO*. Further priority hierarchy can be achieved using the *Select function* of compound models (see definition 2).

Although the library API attempts mimicking the formal structures as closely as possible, there are some fundamental *differences*. For example, the model structures are missing the explicit input and output set parameters at run-time. However, they are present implicitly as the parameter of δ_{ext} , and the output of λ respectively (see definition 1). Furthermore, the input and output parameters are also expressed as *template parameters*. In fact, almost the entire library relies on *template meta-programming*. This approach allows the inputs, outputs, and states to be of *almost any* built-in, or user-defined type. While there are *no requirements* for the input and output types, the *state type* needs to be *copy-able* (due to the *delta* functions), and *printable* using the *overloaded bit-shift operator* «. Another *missing run-time* parameter is the set of influencers I_d from definition 2. Instead, the influencers are *deduced* from the set of *influencer transform functions*. Furthermore, when defining the *influencer transforms*, the *model name*, or *transformer function* can be omitted (by default constructing an optional: {}). If the model name is omitted, the encapsulating *compound model* N is referenced (this way, the components do not need to know the compound model name), and if the transformer is omitted, an *identity function* is used (i.e. value is not changed). An additional difference is that the *Select function* (see definition 2) receives a *vector of model names* instead of a set, making the FIFO selector implementation trivial. This *modified Select* may be defined as $\text{Select}': (2^D, <) \rightarrow D$, where $<$ is an *ordering relation*.

One of the prominent challenges of the library implementation is the *arbitrary connection* and *nesting* of models due to the nature of C++ itself. Both problems are solved by applying *indirection*. The arbitrary model connections could easily be solved by passing around a *generic pointer* that is cast to the desired type by the consumer. The issue with approach is that if the type cast is invalid, the behavior is *undefined*. Instead of this approach, the custom *Dynamic* class is used, which is essentially a thin wrapper around *smart-pointers*. The core idea is using C++'s *dynamic cast* instead of a *raw pointer reinterpretation* due to its ability to *detect and report* invalid casts. This keeps the behavior *defined* even if the library user makes a *mistake* in the model connections, and the user can be properly notified. Thanks to *implicit conversions*, and *operator overloading*, the use of the *Dynamic* class is almost syntactically free (in some cases an explicit typing may be required due to deduction failing/being incorrect). The issue of *model nesting* is solved by defining a model interface named *IOModel* that defines common operations, e.g., *event handling*, *I/O*, etc., and then providing *distinct implementations* for atomic and compound models.

As the project is limited to using only the standard library, only a *CLI* application is feasible. Due to this, an important additional feature is provided in the library, namely *printing*, located in the *Printer* name space. Printers can hook into various events and inform the user. These events include *simulation start*, *simulation step*, *simulation end*, *time advanced*, *event scheduled*, *event action about to be executed*, and finally *model state changed*. Three printers are provided in the library, namely a *base printer* (a *mute printer* that also defines the printer interface, useful for models with unreasonable amount of output), a *verbose printer* (without any text decoration, useful for printing into a file), and finally a *decorated verbose printer* (suitable for printing into the *CLI* directly). The text decorations are based on *ANSI escape sequences*, and the library provides *simple abstractions* around them.

Additional *minor features* include opinionated abstractions for *random number generation* with *uniform*, *exponential*, and *Poisson* distributions. Next, *helper constants* and *types* are defined, e.g., *infinity*, or *null type* (empty class without any special meaning within the library).

The library *API* aims to be as simple and minimal as possible. This means that many structures support *brace-initialization*, and utilize *implicit conversions*, or *default values* wherever appropriate. Unfortunately, parameter deduction does not work correctly in every situation, and explicit parameters may be required.

3 Demo application and provided examples

The project contains an executable *demo application* with six examples built using the created library. Two examples are *minimal* or *empty* models (one atomic, one compound) that showcase the *minimal code* required for running a simulation. Furthermore, these models may be used a foundation for creating complex examples. Next, there is a *traffic light* example simulation, the main purpose of which was testing and developing the library. It is an atomic model with *structured state*, *six input messages*, (power and mode setting/toggling), and *one output message* (possibly the color being switched to).

The final three examples are from the area of *queue theory*. All three examples use the *same queue system model* (depicted

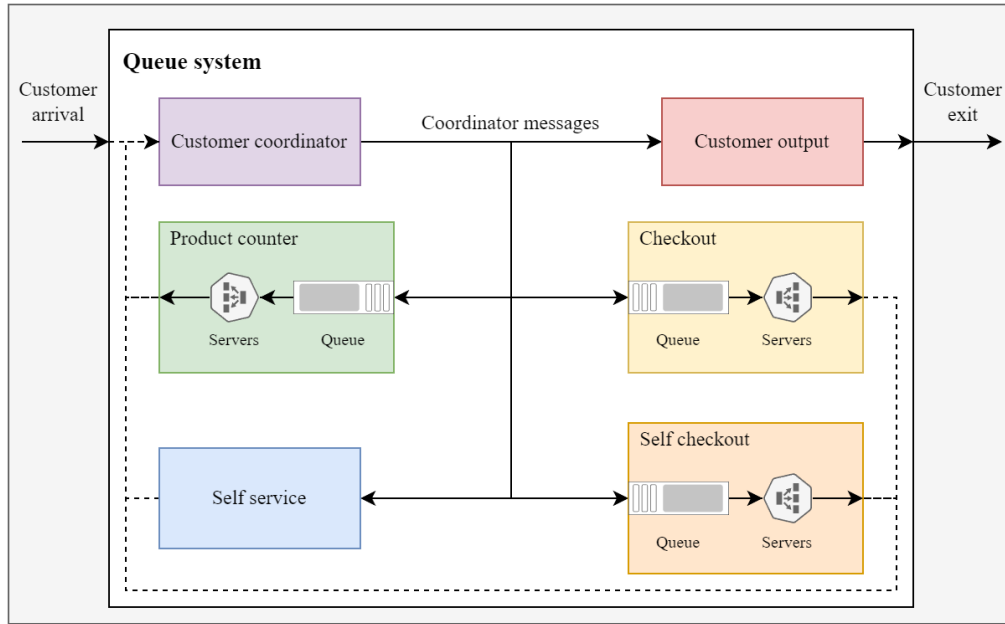


Figure 1. Block diagram representing a queue system from the provided examples (see section 3).

in figure 1) with *similar parameters*. Before any further description, it should be noted that the later referenced *product counter* is not a system counting products, but a *service station*, where customers may request certain products from servers. The queue system is a *compound model* of a grocery store with *four different customers* (requests), *four different service stations*, one type of *output message*, built using *six atomic components*.

The incoming customers of the queue system have *two independent requirements* (yielding *four unique combinations*), namely *product counter* (customers may need to be served at the product counter station), and *age verification* (customers may need to have their age verified at checkout for certain products, e.g., *alcoholic beverages*, possibly increasing the service time). The customer arrivals are scheduled *before the simulation* with *exponentially distributed durations* between arrivals. The flow of customers is defined by the *customer coordinator* component that receives every *new customer*, and customers served by one of the stations. First, the customer is sent to the *product counter*, if *requested*. Next, every customer unconditionally enters a special *self-service* station, representing customers picking products from shelves. The *self-service* station is an extra station *without a queue*, and with *virtually infinite servers* (the customer is its own server, service starts immediately). After self-service, customers proceed to one of the *checkout stations*, based on the *current state* of their respective queues. Before a customer would enter one of the checkouts, the coordinator component sends a *query message* to the checkout stations, and upon *receiving responses* from both, the customer is sent to the station with a *smaller queue*, or always the *checkout* if the queues are *equally large*. Finally after checkout, the customer is removed from the queue system by the *customer output* component. The purpose of this component is *filtering*, *unwrapping*, and *outputting* the customer from the *coordinator message*.

While the *product counter* station is a simple component without additional behaviour, *checkout*, and *self-checkout* are more complex. Specifically, checkouts handle the *age verification* request of customers (each slightly *differently*), and checkouts may have errors. The *advantages* of classic checkouts are a considerably *higher service rate*, *lower error probability*, and that *age verification* does not increase service time (the worker is already there, verification part of the normal service rate). The possible *disadvantages* of classic checkouts are *costs* (building, maintaining), requiring a *human worker* for each server, *footprint*, *lower error handling rate*, and others. On the other hand, *self-checkouts* are more *space efficient*, *cost efficient*, *human workforce efficient*, have *faster error handling* (they are better prepared for error handling as more errors are expected due to increased customer interaction). The downsides of *self-checkouts* include significantly increased *error chance*, *lower service rate*, and that *age verification delays* the service (the human worker may not always be available).

The *purpose* of the queue system model and simulation could be finding the ideal number of servers for every station, or investigating the *ratio of checkout*, and *self-checkout* servers, due to their differing properties, and requirements. However, the demo application contains only fictional parameters, therefore the experiments are tailored towards *differences between variants*, *expectation of behaviour*, or *simulation run-time performance*. As mentioned previously, there are *three* different but similar *parameterizations* of this model in the demo application, specifically:

1. Short duration - In this case, the simulation lasts *ten minutes*, and uses the *colored verbose printer* (additionally, *exiting*

customers print the time of exit), thus writing the output into a text file is not suitable. During this time, around 15-20 customers are expected to arrive. This model is meant to showcase the *step-by-step execution* of components, and present issues of simulations with a low number of customers.

2. Long duration - The parameters of this simulation are identical to the short variant, except instead of 10 minutes, the simulation lasts *ten days*, i.e., a *scaling factor* of 1440.
3. Large scale - The simulation windows for this variant is *twenty-four hours*. Furthermore, the *arrival rate* of customers, and the *number of servers* at every station from the *short variant parameters* is increased by a factor of *ten*, again yielding a total *scaling factor* of 1440.

Running any of the queue variants prints statistics for every station (except self-service). Simulating the short variant yields vastly *differing results*. The *execution time*¹ is anywhere between 500 milliseconds, and 5 seconds. The issue with this simulation is the *time window*. Specifically, the *time window* is *insufficient* in relation to the customer *arrival rate*. Only about 15-20 customers arrive in this simulation, which is not enough to *overcome* the *randomness* of the simulation, contributing to the *instability of results*. Furthermore, the time windows is *too short* in relation to station *service rates*. Often, customers remain *unserved* in one of the stations (although the *total busy times* do reflect pending customers, meaning that busy times may sometimes be over 100%, indicating that the server had to work overtime). Another negative factor is the *amount of text output*, further slowing down the execution. Next, the *long variant* has the same parameters as the short, except the time window is *ten days*. The *execution time* of this simulation is consistently close to 3.5 seconds. During this time, around 24 000 customers are served in the system, which is almost exactly 1440 times more than in the short variant, indicating *correct scaling*. Due to the large number of *served customers*, the randomness becomes insignificant, yielding more *stable results*. Finally, the *large variant* yields *comparable results* to the long variant, with one *downside*. Even though the total scaling factor for both is 1440 in relation to the short, the execution time for the large variant is consistently over 8 seconds, i.e., more than double of the long variant. This is likely due to the *larger number of servers*, noticeably slowing down state updates.

In conclusion, results of the *short variant* are unpredictable due to a short time window. The only significant *difference* between the long and large variants is the *execution time*, with the large variant requiring more than *twice* the time of the long variant, making the *long variant* the best choice, given the fictional parameters, and implementation details. Both long and large variants yield relatively *stable*, and *predictable* results. For example, the *product counter* has *half the service rate* of *total arrival rate*, compensated by two servers, meaning that on average, the product counter should be able to handle all incoming customers. However, only about 75% of customers request service at the *product counter*. Thus, the product counter should be *busy* for around 75% of the total duration, and *idle* for the remaining 25%. Furthermore, *checkouts* are slightly *busier* than *self-checkouts*, and the average queue size of checkouts is larger by about 0.5, caused by always selecting the checkouts, when the queue sizes are equal. Simulations of the long and large variants consistently yield the expected results (max. 1% error).

4 Conclusions

Firstly, this technical report introduces the *DEVS* formalism, and its *atomic*, and *compound* models in section 1. Next in section 2, *features* and interesting *implementation details* of the library are discussed. Although the library attempts *mimicking* the *formal definition* as closely as possible, there are minor differences. Another aim of the library is being *general*, covering most *use-cases*, while remaining *unopinionated*. This is achieved using *template meta-programming*, allowing any *input*, *output*, or almost any *state type* (state needs to be *copy-able* due to the definition of δ functions of atomic models, and *printable* using the `<<` operator). This also means it *does not provide* abstractions for these types, making the examples rather lengthy. However designing a library for these types is beyond the scope of this project. Finally, the provided *demo application* and its *examples* are presented in section 3. Three of the six provided examples are from the *queue theory* area, sharing the *same model* with *different parameters*. Variants of the parameterization are referred to as *short* (*short time windows*), *long* (very long time window), and *large* (moderate time window, scaled arrival rates, scaled servers, implying more customers present in the system on average). *Experiments* indicate that the best variant is *long*, likely caused by *implementation details*. The project hopefully provides a good *foundation* and *potential* for additional work, e.g., *saving/loading* of models into/from files, *improved visualization*, or *acceleration using threads*.

References

- [1] Petr Peringer. “SNT — Simulační nástroje a techniky”. Brno, 2023.
- [2] Bernard P. Zeigler, Alexandre Muzy, and Ernesto Kofman. *Theory of modeling and simulation: discrete event and iterative system computational foundations*. 3rd ed. London: Academic Press, 2019. ISBN: 978-0-12-813370-5.

¹The provided execution times are from running on a laptop with an i9-12900H CPU, Windows 11 x64, in a WSL environment running Ubuntu 20.04.4.